

# DEEP LEARNING INFERENCE USING INTEL<sup>®</sup> OPENVINO<sup>™</sup> TOOLKIT

Elroy Ashtian, Jr.

11/7/2018

Core and Visual Computing Group

OpenVINO toolkit overview video: Innovate Vision Solutions with OpenVINO™ toolkit



0:00 / 2:53





**EMERGENCY RESPONSE**



**FINANCIAL SERVICES**



**MACHINE VISION**



**CITIES/TRANSPORTATION**

# VIDEO: THE “EYE OF IOT”

USE OF VIDEO, COMPUTER VISION AND DEEP LEARNING IS GROWING RAPIDLY



**AUTONOMOUS VEHICLES**



**RESPONSIVE RETAIL**



**MANUFACTURING**

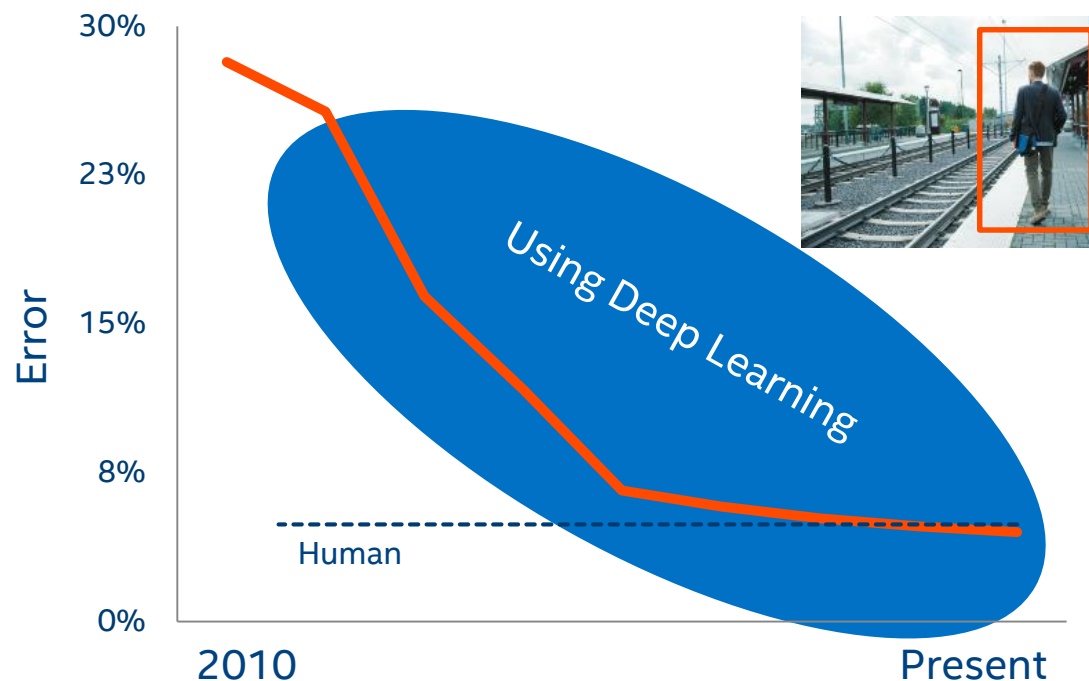


**PUBLIC SECTOR**

# Deep Learning Usage Is Increasing

Deep learning revenue is estimated to grow from \$655M in 2016 to **\$35B** by 2025<sup>1</sup>.

## Image Recognition



Traditional  
Computer Vision  
Object Detection



Deep Learning  
Computer Vision  
Person  
Recognition




Market Opportunities + Advanced Technologies Have Accelerated Deep Learning Adoption


<sup>1</sup>Tractica\* 2Q 2017

# Deep Learning: Training vs. Inference

## Training




## Inference



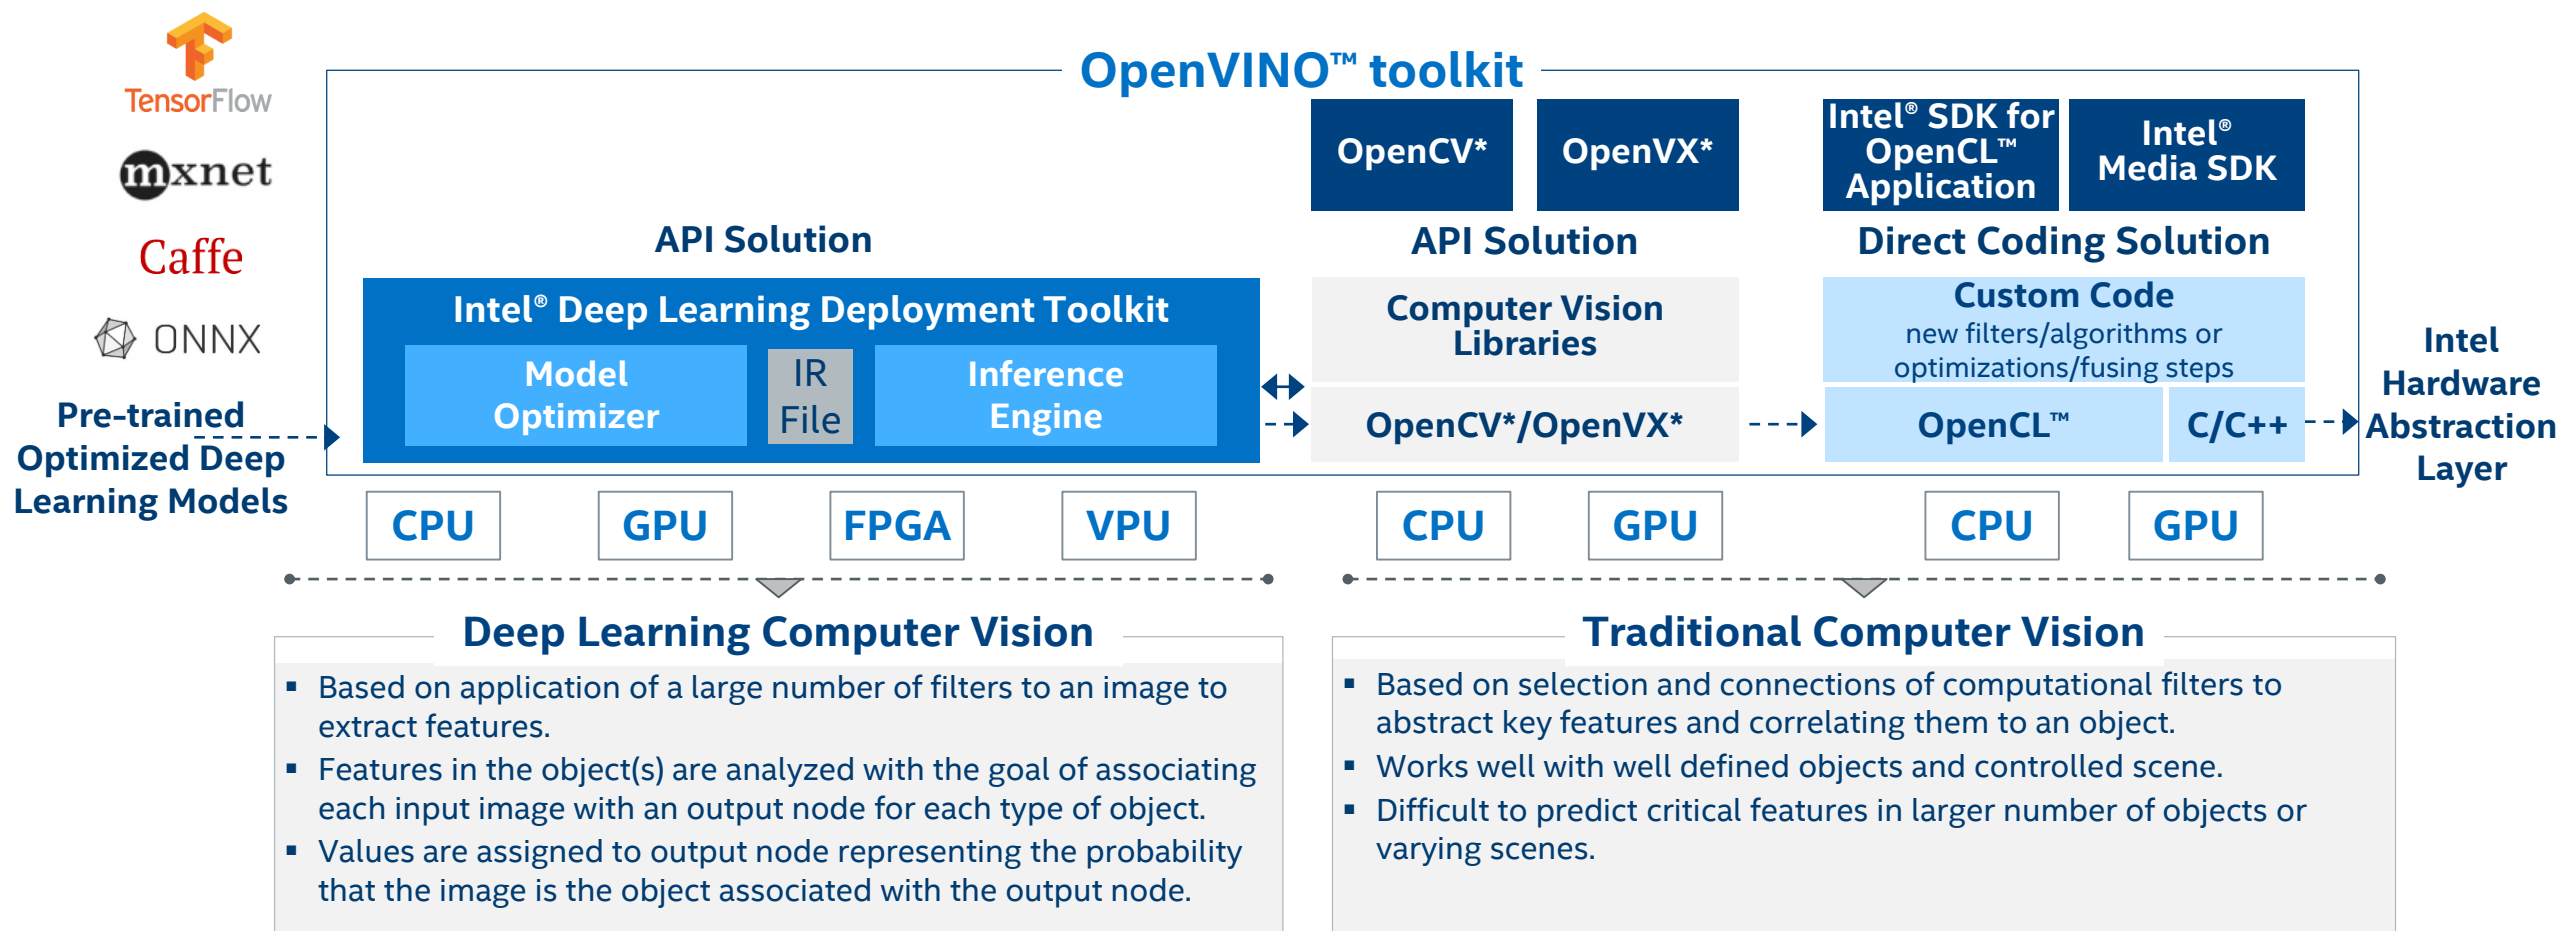
## Did You Know?

Training requires a very large data set and deep neural network (many layers) to achieve the highest accuracy in most cases





# Open Visual Inference & Neural network Optimization (OpenVINO™) toolkit



IR = Intermediate Representation File  
GPU = Intel CPU with integrated graphics processing unit/Intel® Processor Graphics  
VPU = Intel® Movidius™ Vision Processing Unit

## Optimization Notice

Copyright © 2018, Intel Corporation. All rights reserved.  
\*Other names and brands may be claimed as the property of others.

OpenVX and the OpenVX logo are trademarks of the Khronos Group Inc.  
OpenCL and the OpenCL logo are trademarks of Apple Inc. used by permission by Khronos



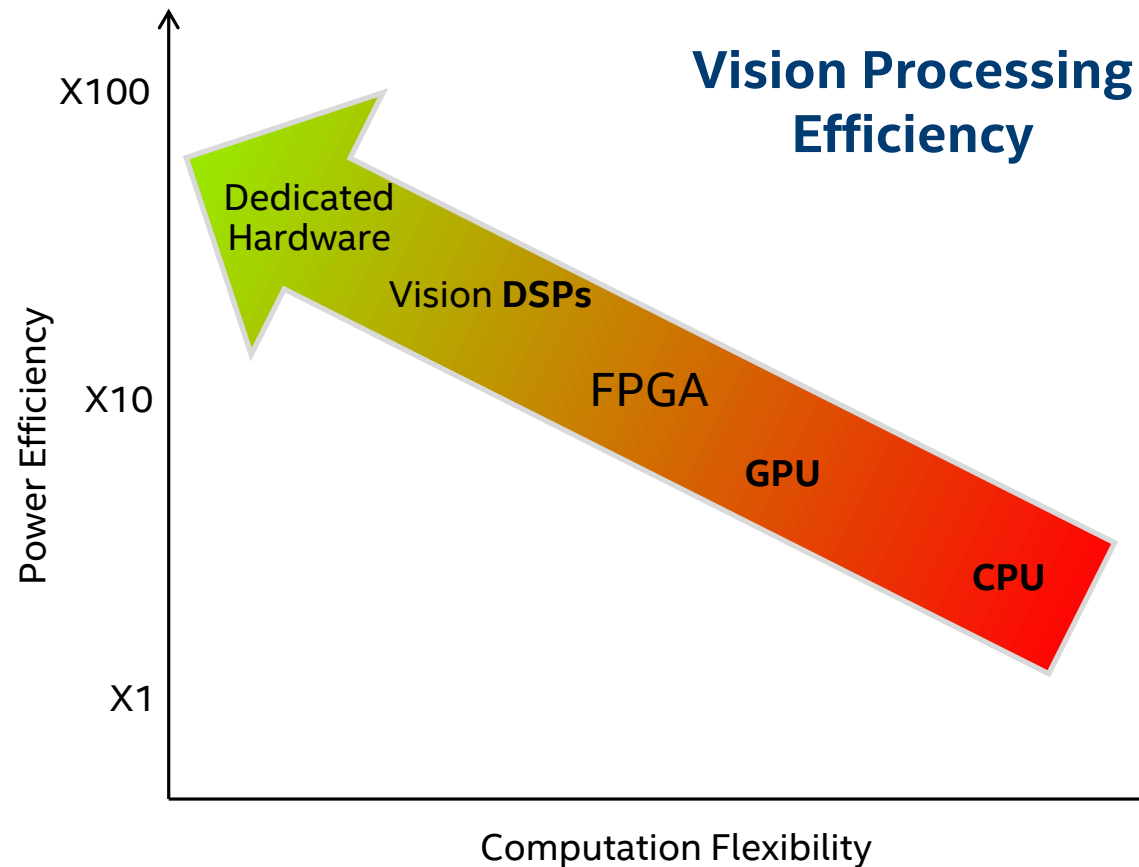
# Choosing the “Right” Hardware

## Considerations

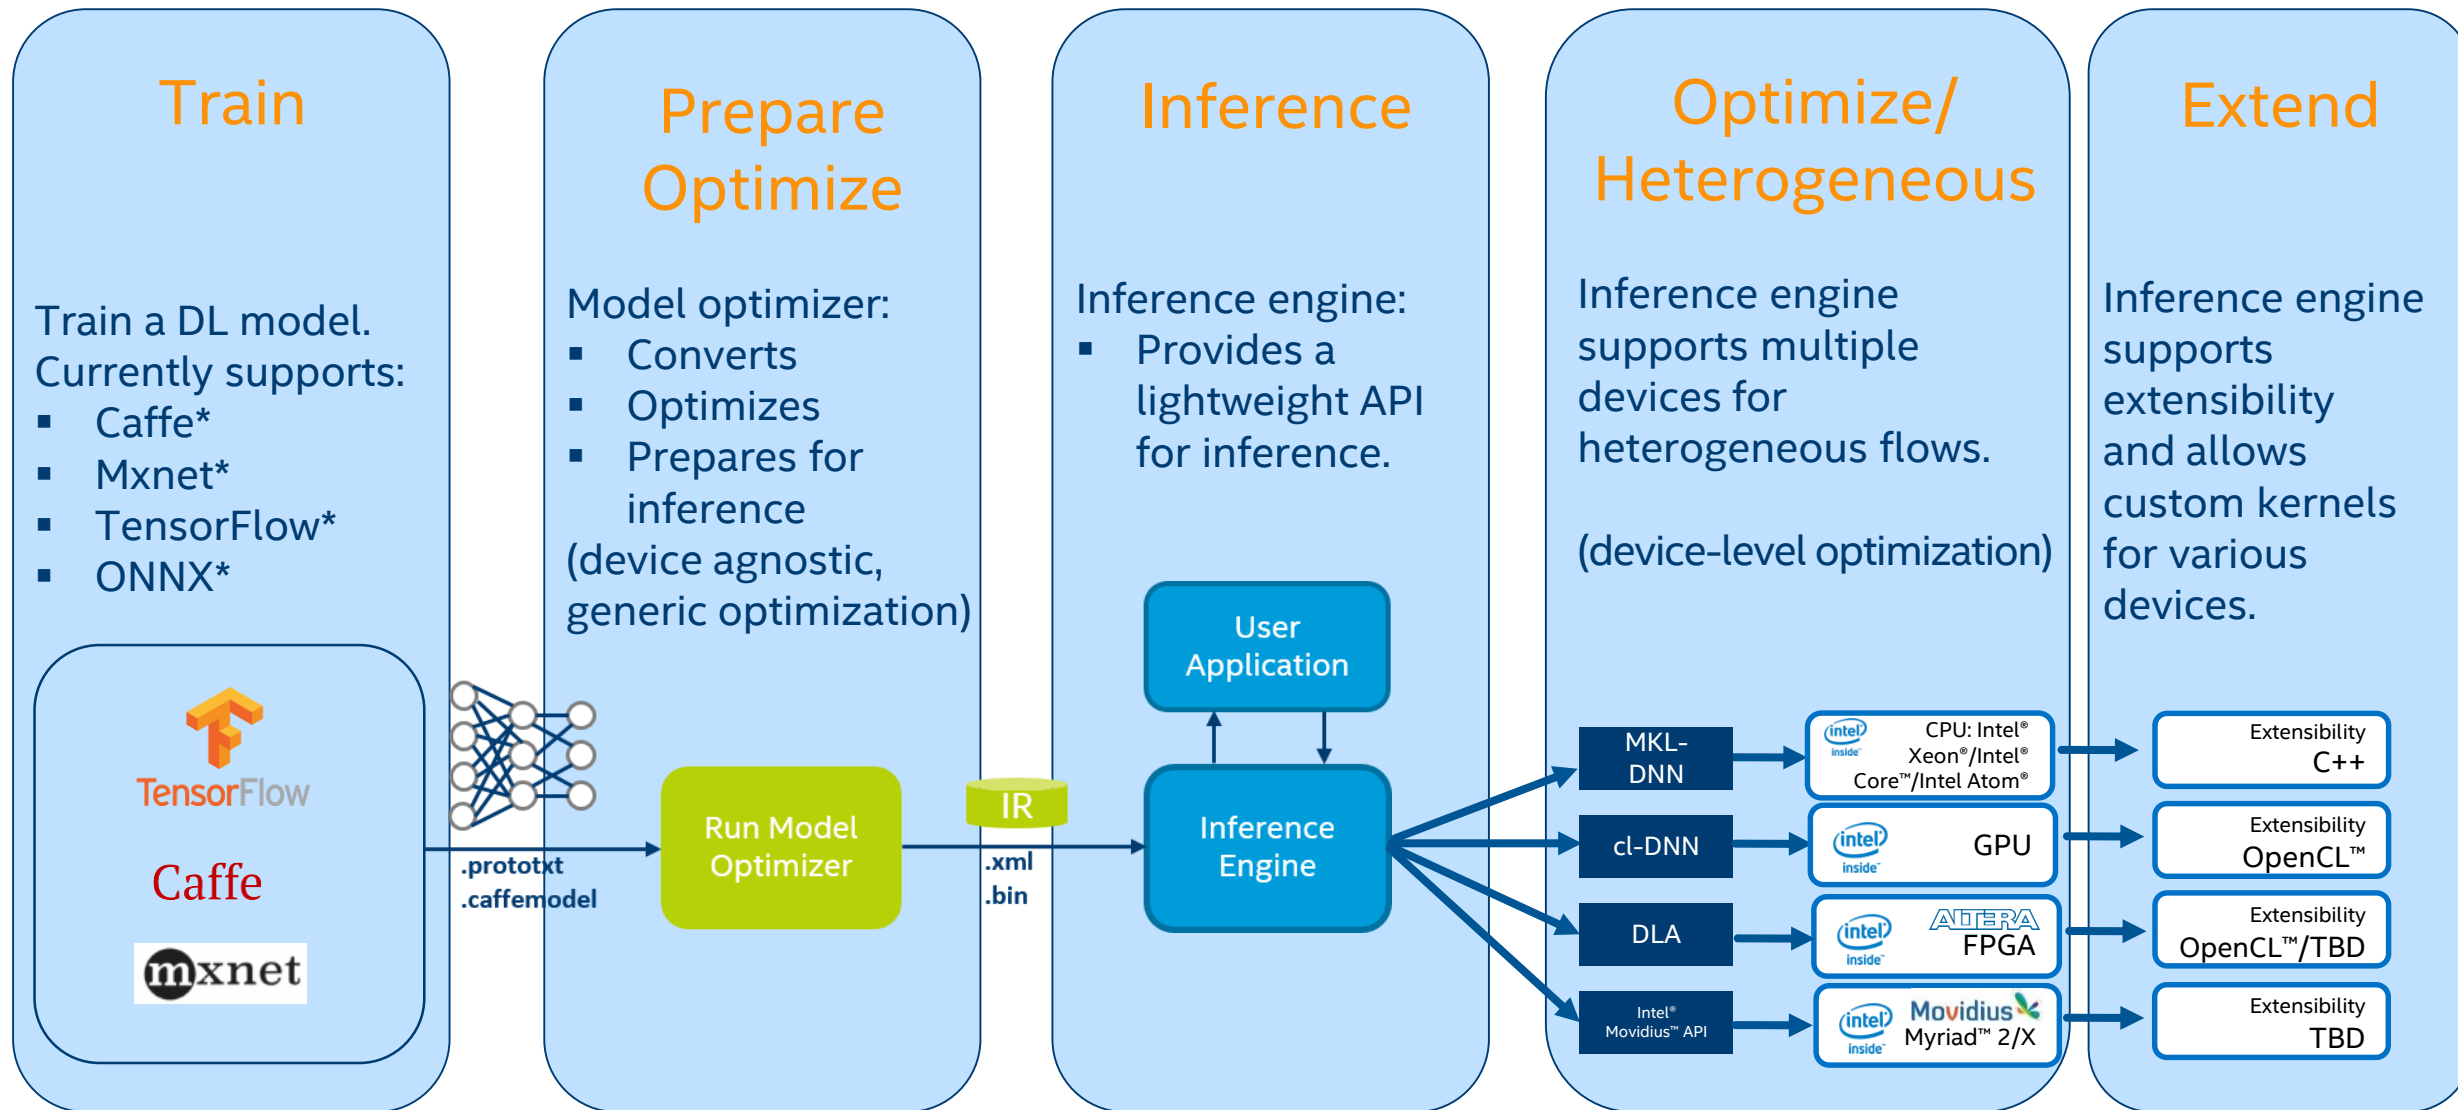
- Workload
- Hardware acceleration
- Heterogeneous computing

## Tradeoffs

- Power/performance
- Price
- Software flexibility, portability




# Application development with OpenVINO™ Toolkit





# Deep Learning Workload Performance on Public Models using OpenVINO™ toolkit & Intel® Architecture

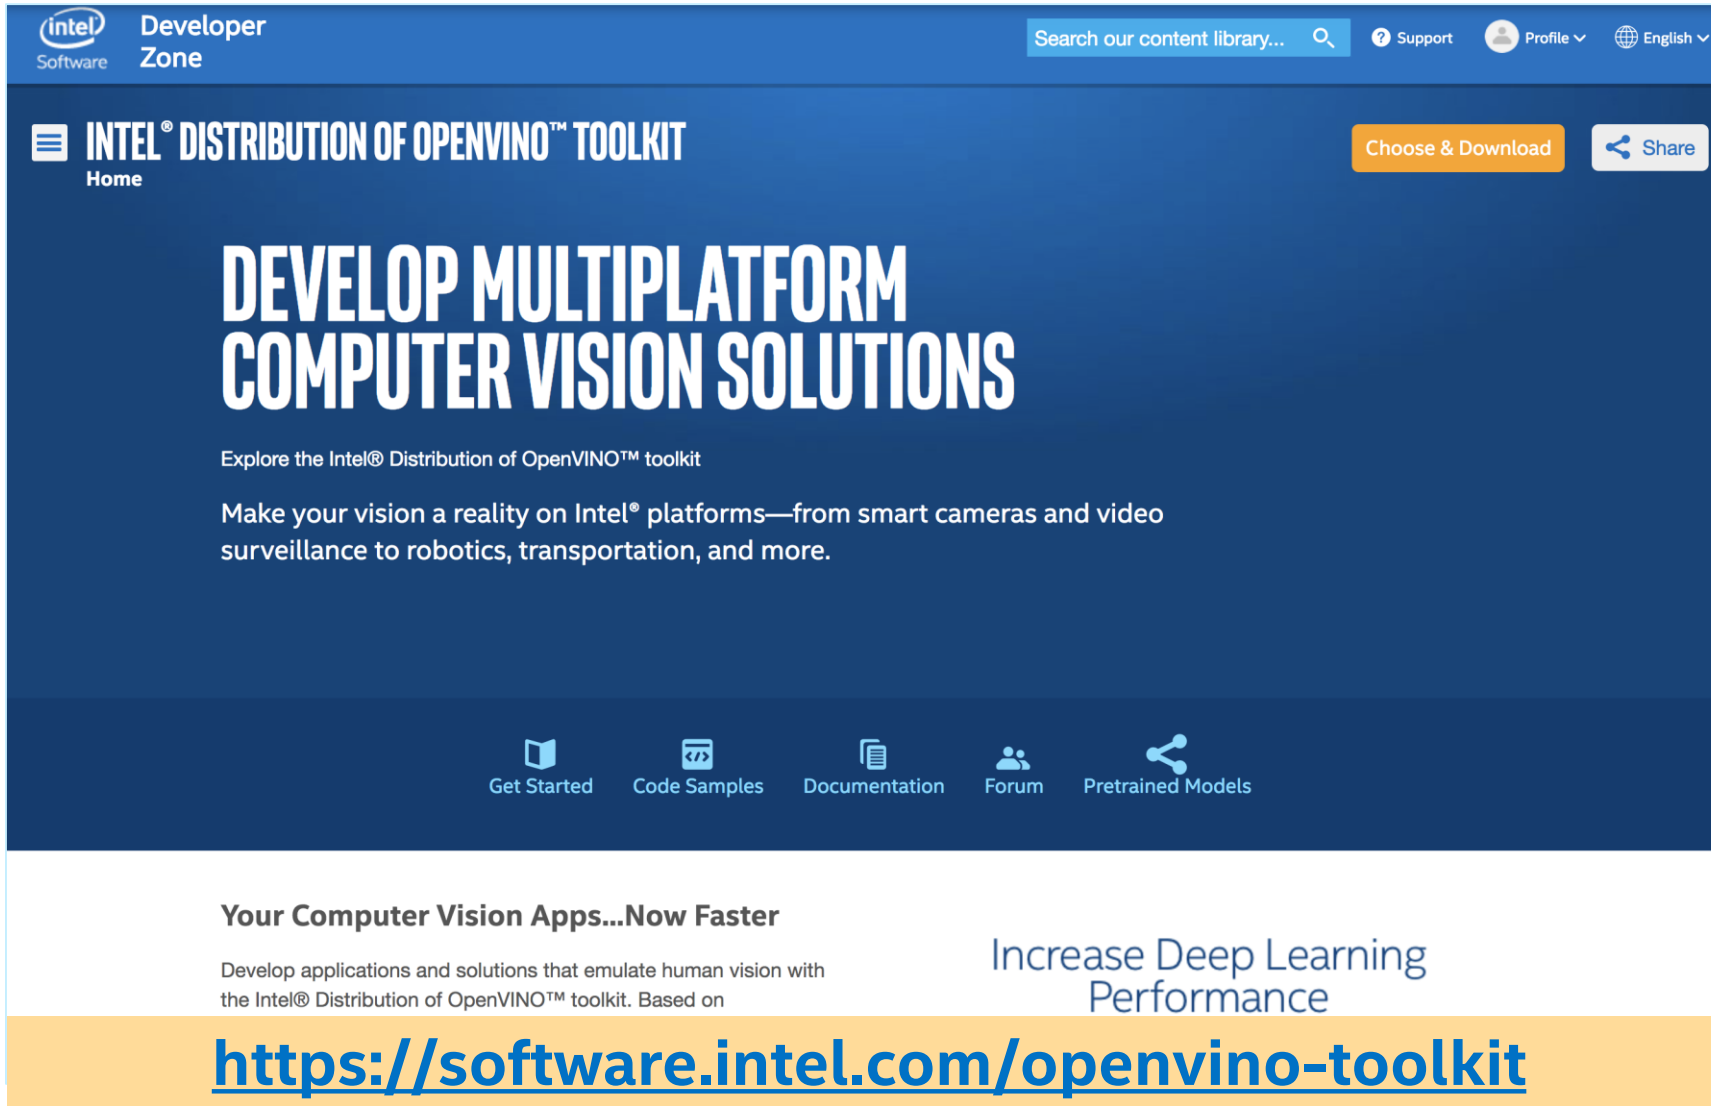
## Comparison of Frames per Second (FPS)



Get an even Bigger Performance Boost with Intel® FPGA

<sup>1</sup>Depending on workload, quality/resolution for FP16 may be marginally impacted. A performance/quality tradeoff from FP32 to FP16 can affect accuracy; customers are encouraged to experiment to find what works best for their situation. Performance results are based on testing as of June 13, 2018 and may not reflect all publicly available security updates. See configuration disclosure for details. No product can be absolutely secure. For more complete information about performance and benchmark results, visit [www.intel.com/benchmarks](http://www.intel.com/benchmarks). **Configuration:** Testing by Intel as of June 13, 2018. Intel® Core™ i7-6700K CPU @ 2.90GHz fixed, GPU GT2 @ 1.00GHz fixed Internal ONLY testing, Test v3.15.21 – Ubuntu\* 16.04, OpenVINO 2018 RC4, Intel® Arria® 10 FPGA 1150GX. Tests were based on various parameters such as model used (these are public), batch size, and other factors. Different models can be accelerated with different Intel hardware solutions, yet use the same Intel software tools.

# Intel® OpenVINO™ toolkit on Intel Developer Zone



The screenshot shows the Intel Developer Zone website for the OpenVINO toolkit. The header includes the Intel Software logo, 'Developer Zone' text, a search bar, and links for Support, Profile, and Language. The main content area features a large heading 'DEVELOP MULTIPLATFORM COMPUTER VISION SOLUTIONS' and a subheading 'EXPLORE THE INTEL® DISTRIBUTION OF OPENVINO™ TOOLKIT'. Below this, a paragraph states: 'Make your vision a reality on Intel® platforms—from smart cameras and video surveillance to robotics, transportation, and more.' A navigation bar contains icons and labels for 'Get Started', 'Code Samples', 'Documentation', 'Forum', and 'Pretrained Models'. At the bottom, there are two promotional sections: 'Your Computer Vision Apps...Now Faster' and 'Increase Deep Learning Performance'. A large orange banner at the very bottom displays the URL <https://software.intel.com/openvino-toolkit>.

intel Software Developer Zone

Search our content library...

Support Profile English

INTEL® DISTRIBUTION OF OPENVINO™ TOOLKIT

Home

Choose & Download Share

## DEVELOP MULTIPLATFORM COMPUTER VISION SOLUTIONS

Explore the Intel® Distribution of OpenVINO™ toolkit

Make your vision a reality on Intel® platforms—from smart cameras and video surveillance to robotics, transportation, and more.

Get Started Code Samples Documentation Forum Pretrained Models

**Your Computer Vision Apps...Now Faster**

Develop applications and solutions that emulate human vision with the Intel® Distribution of OpenVINO™ toolkit. Based on


Increase Deep Learning Performance

<https://software.intel.com/openvino-toolkit>

# Inference with Intel® OpenVINO



(a) classification



(b) detection



(c) segmentation

<https://software.intel.com/openvino-toolkit/documentation/code-samples>

# Day 2 – November 8<sup>th</sup> (Hands-on Labs)

09:00 AM - 09:45 AM	Breakfast
10:00 AM - 11:00 AM	Video Performance using Intel® Media SDK
11:00 AM - 12:00 PM	Hardware based Remote Management using Intel® AMT
12:00 PM - 01:00 PM	Lunch
01:00 PM - 05:00 PM	Video Analytics using Intel® OpenVINO™ <ul style="list-style-type: none"><li>• Face Detection using OpenVINO™ toolkit lab</li><li>• Age and Gender Detection using OpenVINO™ toolkit lab</li><li>• Face Data Analysis on the Cloud Lab</li></ul>

